



OSI – M30233 (OS Theme)

Dr Tamer Elboghdadly

Synchronization and Deadlock

- **Plan:**
 - General kinds of synchronization
 - Resource deadlocks, and
 - Strategies for dealing with them
- **There is an “appendix”**
 - Notification synchronization – detailed example

Synchronization

Mutual Exclusion Reprise

- ***ME*** is one of the most important kinds of *synchronization* between threads.
 - No thread can enter a critical section while another thread is inside its own critical section.
 - Can be implemented with the aid of various software devices, e.g. *semaphores*.

Beyond Mutual Exclusion

- Many other kinds of synchronization, e.g.:
 - Thread i *sends a message* to thread j
 - *Receive* operation in thread j can't complete until *send* operation in thread i is reached.
 - Sometimes *send* can't complete until thread j reaches *receive* – depends how much buffering.
 - *Join* synchronization between parent and child threads
 - *Join* operation in parent can complete when child terminates.
 - *Barrier* synchronization across a group of *N* processes.
 - No thread can complete its *barrier* operation until all *N* threads have reached their barrier operation.

General Synchronization

- Generally “synchronization” consists in a particular thread having to wait until some condition is created by one or more other threads.
- *Dijkstra’s semaphores*, introduced last week, are one rather general mechanism for achieving different kinds of synchronization.
 - Recall semaphore S is an integer supporting two operations:
 - $P(S)$ decrease semaphore, *but not below zero*
 - $V(S)$ increase semaphore.

Other Uses of Semaphores

- In ME, semaphore is usually initialized to 1.
- In following example we assume instead semaphore S is initialized to 0 :

<i>Thread i</i>	<i>Thread j</i>
\vdots	\vdots
A	$P(S)$
$V(S)$	B

- Execute B in thread j *only after* A executed in thread i .
- *Notification* synchronization – compare with send/receive synchronization.
- *See appendix to this week's lecture for more!*

RESOURCE DEADLOCK

Resources


- Computer systems have many kinds of *resources*, that can only be *accessed by one process or thread at a time*.
- Examples include shared data structures in the operating system (mutual exclusion), and physical devices such as printers.
 - If two threads update a shared data structure at the same time, likely result is corruption of the data structure.
 - If two processes send data to a printer at the same time, likely output is gibberish.

Deadlocks

- *Resource deadlocks* can occur when processes (or threads) need to acquire access to more than one exclusive resource.
 - For example, a programme may need to use a scanner and a printer, and it acquires exclusive access to both these devices.
- Classic example involves two processes *A* and *B* and two resources *R* and *S*.
 - Both processes want to use both resources, but try to acquire them in different orders.

Deadlock Examples

- Example 1:
 - System has 1 scanner and 1 printer.
 - Processes *A* and *B* each hold one each of scanner and printer, and each needs the other one.



Deadlock Examples

- Example 2:
 - Semaphores R and S , initialized to 1.

<i>Thread i</i>	<i>Thread j</i>
-----------------	-----------------

·	·
·	·
·	·

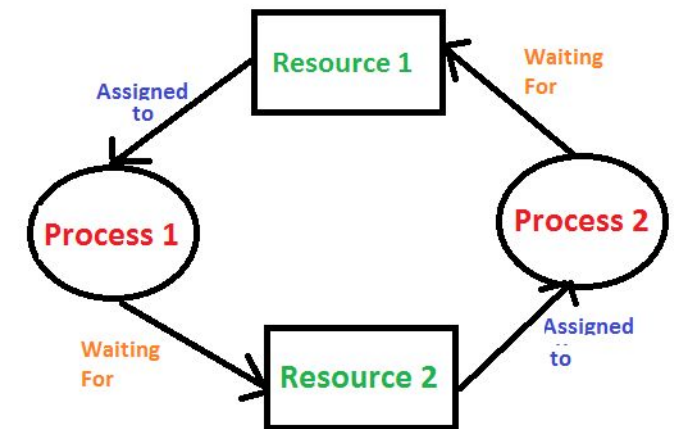
$P(R)$	$P(S)$
--------	--------

$P(S)$	$P(R)$
--------	--------

·	·
·	·
·	·

What is Deadlock?

- Deadlock is a situation where a process or a set of processes wait indefinitely for an event that can never occur.
- More specifically:
 - *A set of threads is in a **resource deadlock** state when every thread in the set is waiting on a resource which is being held by another thread in the set.*





Deadlock Modelling

- Resource deadlock can be modelled using a *Resource Allocation Graph*, which shows:
 - Which processes are requesting which resources.
 - Which resources have been granted to which processes.
- If graph contains no cycles ➡ no deadlock.
- If graph contains a cycle ➡ deadlock.

Resource Allocation Graphs

- Directed graphs[†]:

Note: Directions of the arrows are important.



- resource **R** assigned to process **A**
- process **B** is **requesting/waiting** for resource **S**
- process **C** and **D** are in **deadlock** over resources **T** and **U**

Process:



Resource:



How A Deadlock May Occur†

A
Request R
Request S
Release R
Release S
(a)


B
Request S
Request T
Release S
Release T
(b)

C
Request T
Request R
Release T
Release R
(c)


No Deadlock

1. A requests R
2. B requests S
3. C requests T
4. A requests S
5. B requests T
6. C requests R


(d)



(e)



(f)




(g)

Deadlock


Process:




Resource:



(h)



(i)



(j)

†Tanenbaum, MOS, Fig 6-4

Slide 16 of 32

Thread B acquires lock
Thread A waiting
Thread C waiting

Thread B acquires lock
Thread A waiting
Thread C waiting

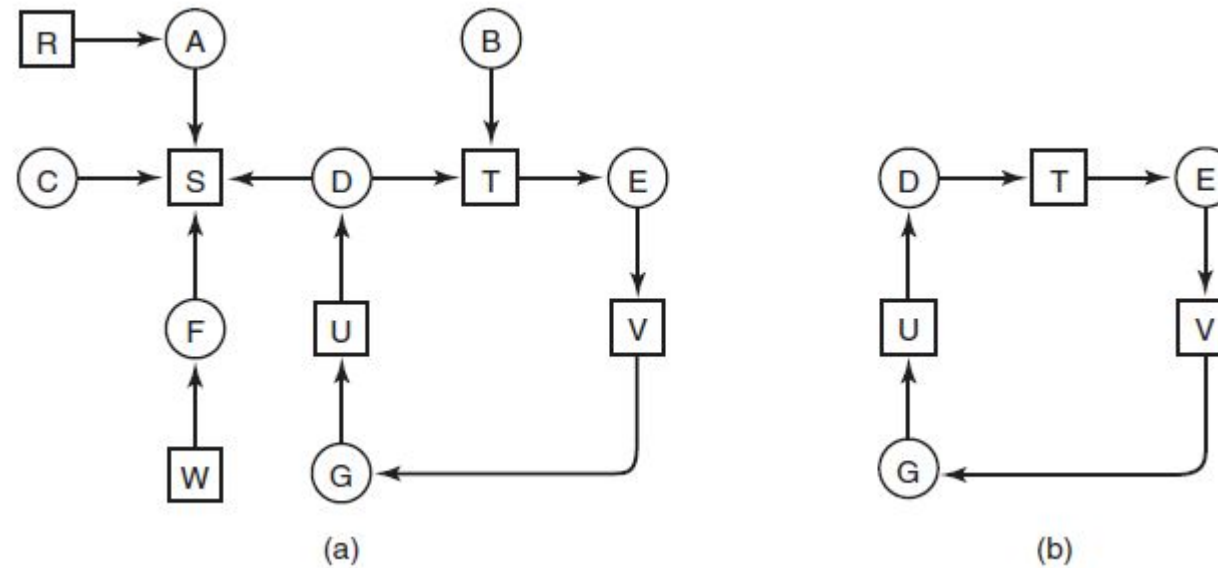


DEALING WITH DEADLOCK

Deadlock Detection and Recovery

- **Basic idea:**
 - Allows system to enter deadlock state.
 - Runs Detection algorithm periodically to check.
 - Performs Recovery scheme if deadlock.
- **To detect deadlock, search for a cycles in the Resource Allocation Graph.**

Deadlock Detection†



- The graph **a)** contains the cycle **b)**.
- This indicates **deadlock**

†Tanenbaum, MOS, Fig 6-5


Deadlock Recovery

- Crudely - kill processes until the deadlock cycle is eliminated.
- “Survivors” get the resources.
- Commonly used in *Relational Database Management Systems*, where transactions identified as causing deadlock can be “rolled back”.

Deadlock Avoidance

- The idea here is to keep the system “safe” – avoid entering “unsafe” states which may later turn into a deadlock.
- Best illustrated by an example.
 - Consider the classic deadlock scenario of two processes sharing resources
 - Each process has its own “instruction counter”.
 - If we plot these as two axes of a graph, behaviour of whole system is represented as a trajectory in this graph, moving **upwards** and to the **right**.

Two Processes and one Resource



Remarks

- Both processes request R , then later release it.
- The **green** trajectory is a possible evolution of the system as a whole – it is “moving” upwards and to the right (**can never go down or to the left**, because that implies process A or B is “going backwards”).
- **Blue** region is forbidden by mutual exclusion (note A reaches this critical section last, and must wait).

Two Processes and two Resources

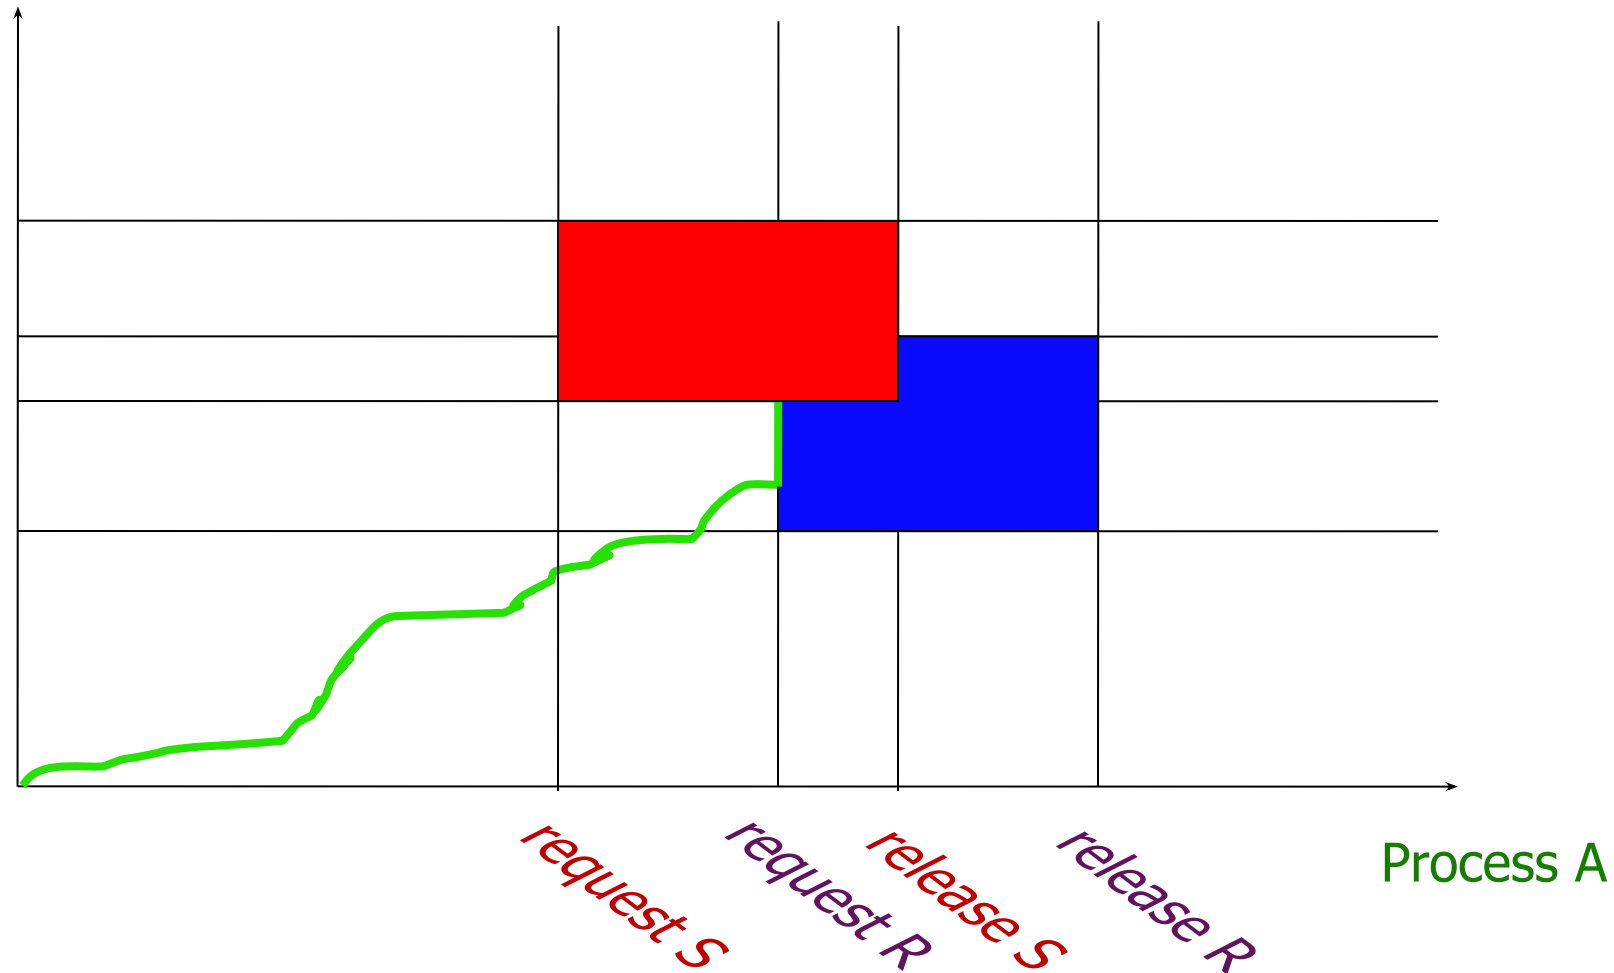
Process B

release S

release R

request S

request R



Remarks

- Process *A* requests *S* then *R*; Process *B* requests *R* then *S*.
- Blue region is forbidden by mutual exclusion on *R*; Red region forbidden by mutual exclusion on *S*.
- The green trajectory is now blocked when it reaches the intersection of the two forbidden regions
 - system can only evolve *up or to right* - means it can't evolve at all from here!

Unsafe Region


Process B

release S

release R

request S

request R



Process A

Safe and Unsafe Regions

- The amber region in the previous graph is an *unsafe region*.
- If the system *ever* enters this region, it will *inevitably* be driven to the point of deadlock
 - because it can't move “down” or “to the left”.
- A *deadlock avoidance* scheme should delay acquisition of any resource, if acquiring it would allow the system into an unsafe region.

Banker's Algorithm (Dijkstra, 1965)


- Basic idea:
 - Requires all processes to declare the max # of resource units that they may request.
 - Keeps track of current allocation for each process and their current “needs” (“max” – “has”).
 - When receives a request, “pretend” to honor the request, and
 - Try to fulfill the “needs” of all other processes in some order so as to check whether or not granting the request will lead to a safe state.
 - If safe, grant request; otherwise, deny request.
- [Read Modern Operating Systems, sections 6.5.3 – 6.5.4](#)

Summary

- General synchronization
 - Beyond mutual exclusion, and more on semaphores
 - Notification using monitors
- Resource deadlocks, and how to deal with them
 - Detect and recover (Detection algorithm).
 - Avoid by not entering an unsafe state (Banker's algorithm),
- Next lecture: *Processes and scheduling*

Further Reading

- Andrew S. Tanenbaum, “*Modern Operating Systems*”, 4th Edition, Pearson, 2014 (MOS), Chapter 6.



UNIVERSITY^{OF}
PORTSMOUTH

Questions?

