



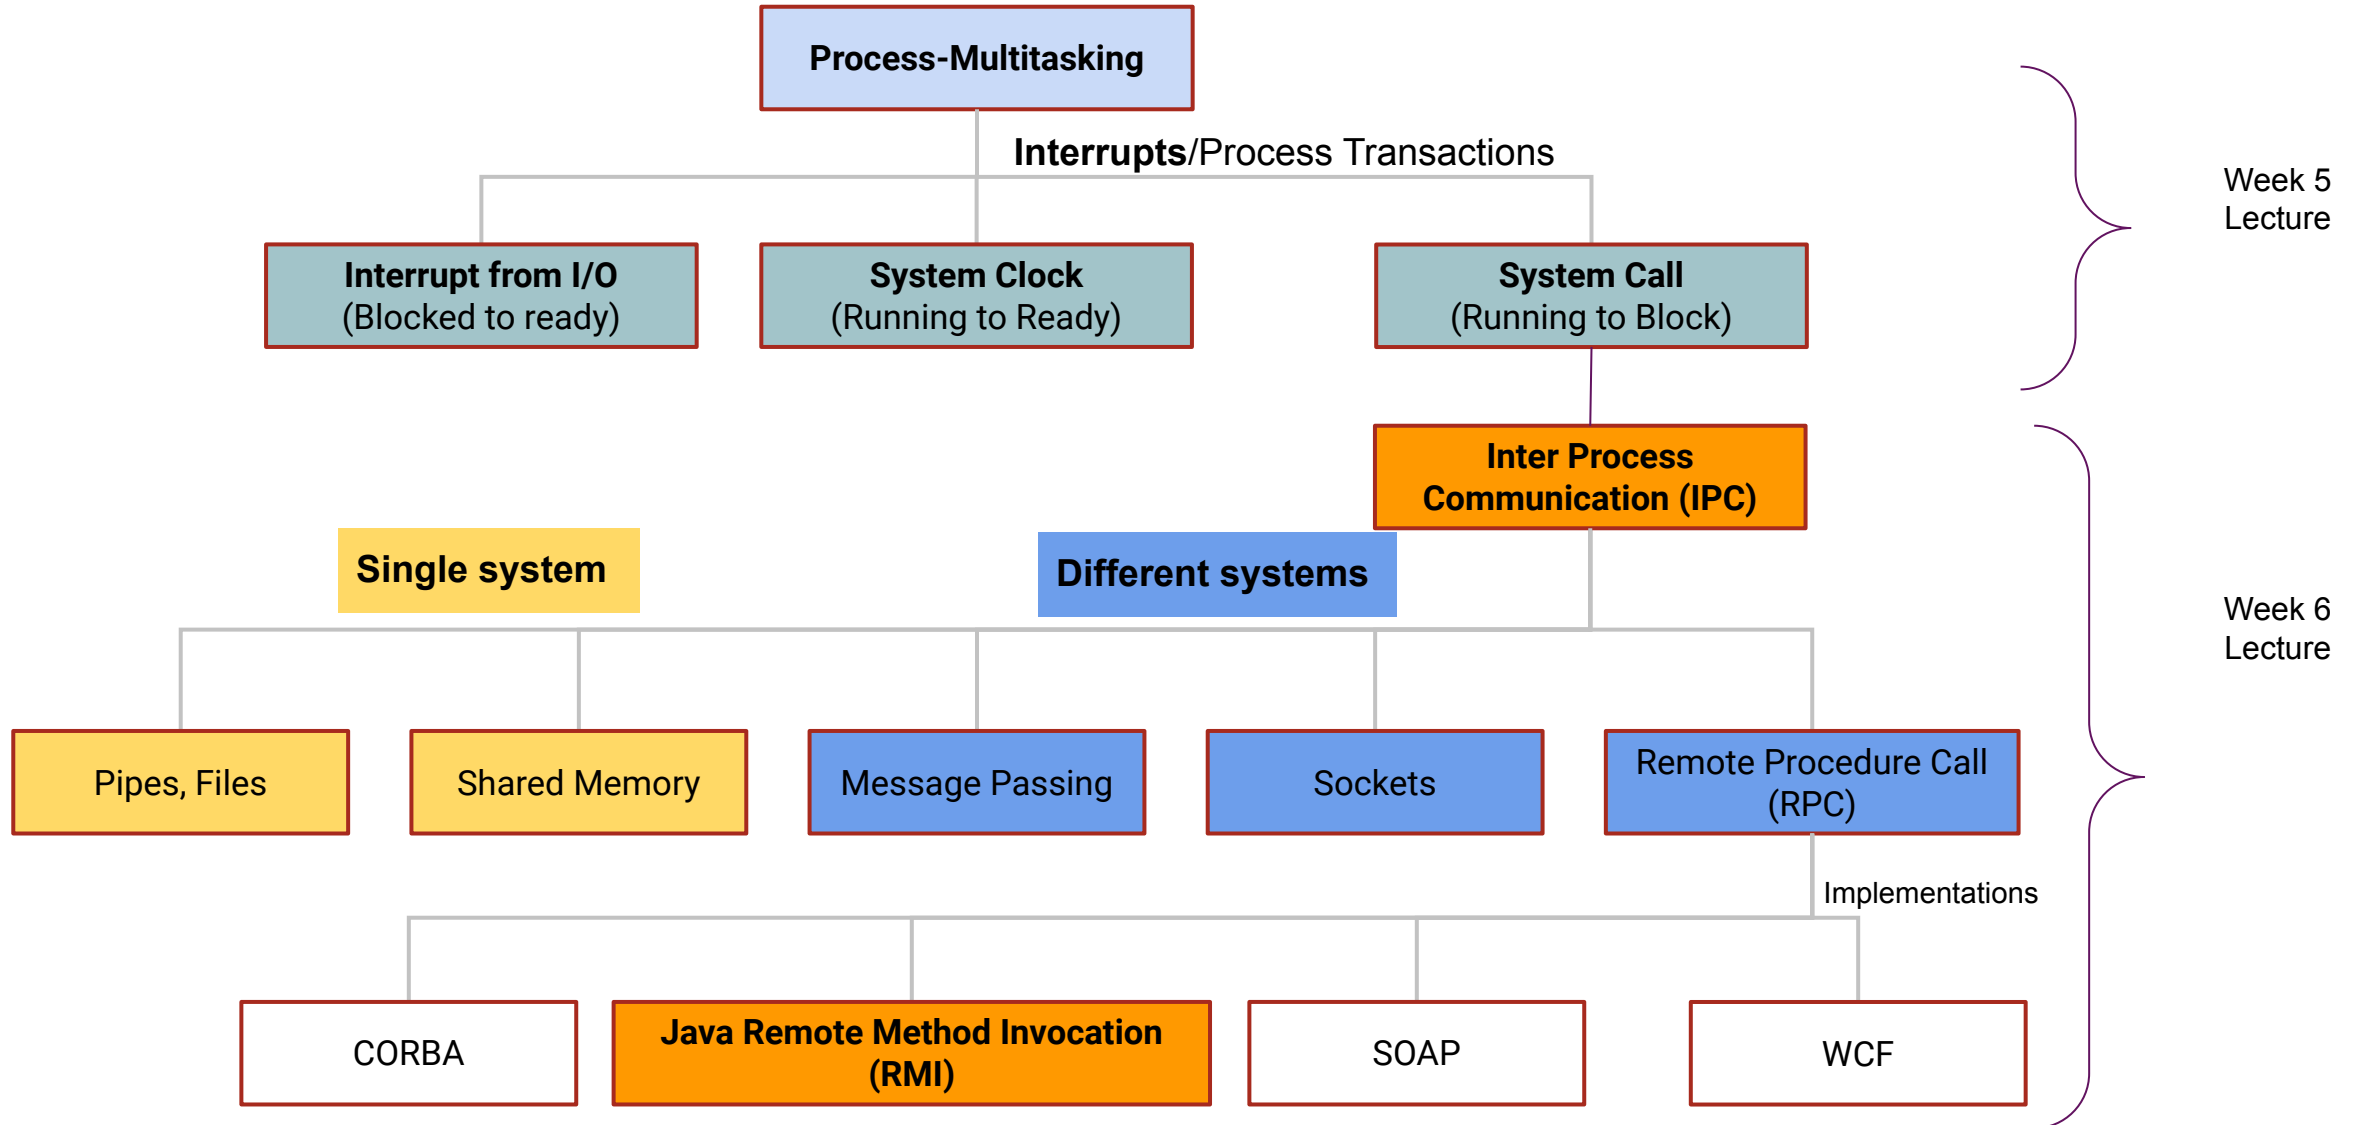
OSI – M30233 (OS Theme)

Dr Tamer Elboghdadly

Interprocess Communication

- **Plan:**

- Interprocess Communication (IPC) overview
- IPC for communication between processes in “**single systems**”
- IPC for communication between processes running on **different systems**.



Inter-process Communication (IPC)

- In earlier lectures we considered interaction between *threads* in general, using access to shared variables, and semaphores, etc.
- But *processes* generally have their own private address space, and don't have any program variables in common.
- Traditionally, OS-supported mechanisms by which processes running in the same computer cooperate are called *Inter-process Communication* (IPC)
- We will take a broader view of IPC, and include mechanisms through which processes on *different* computers can communicate.

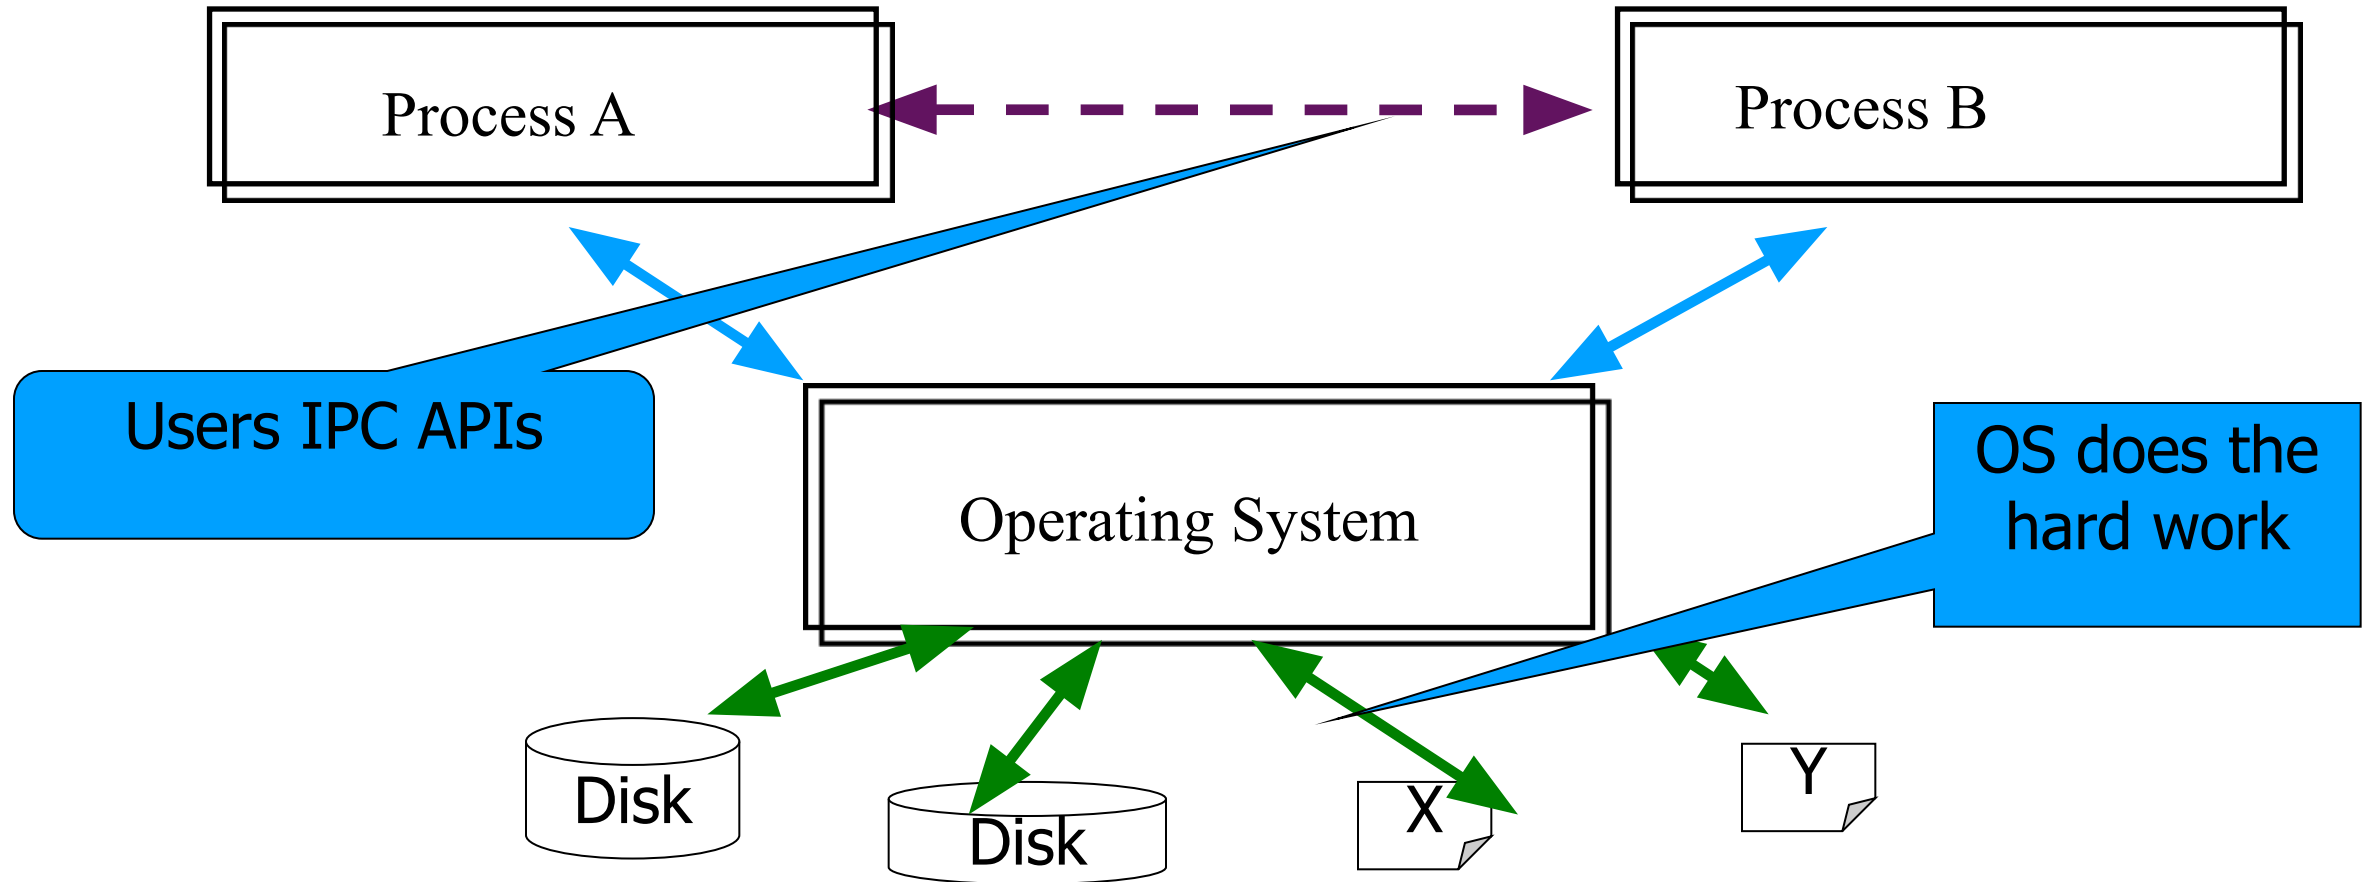
Benefits and Problems

- Advantages of process cooperation:
 - Information sharing,
 - Computation speed-up (parallel processing),
 - Modularity.
- Questions:
 - How do the processes “share” or “communicate” data?
 - What are the problems/issues of sharing data?

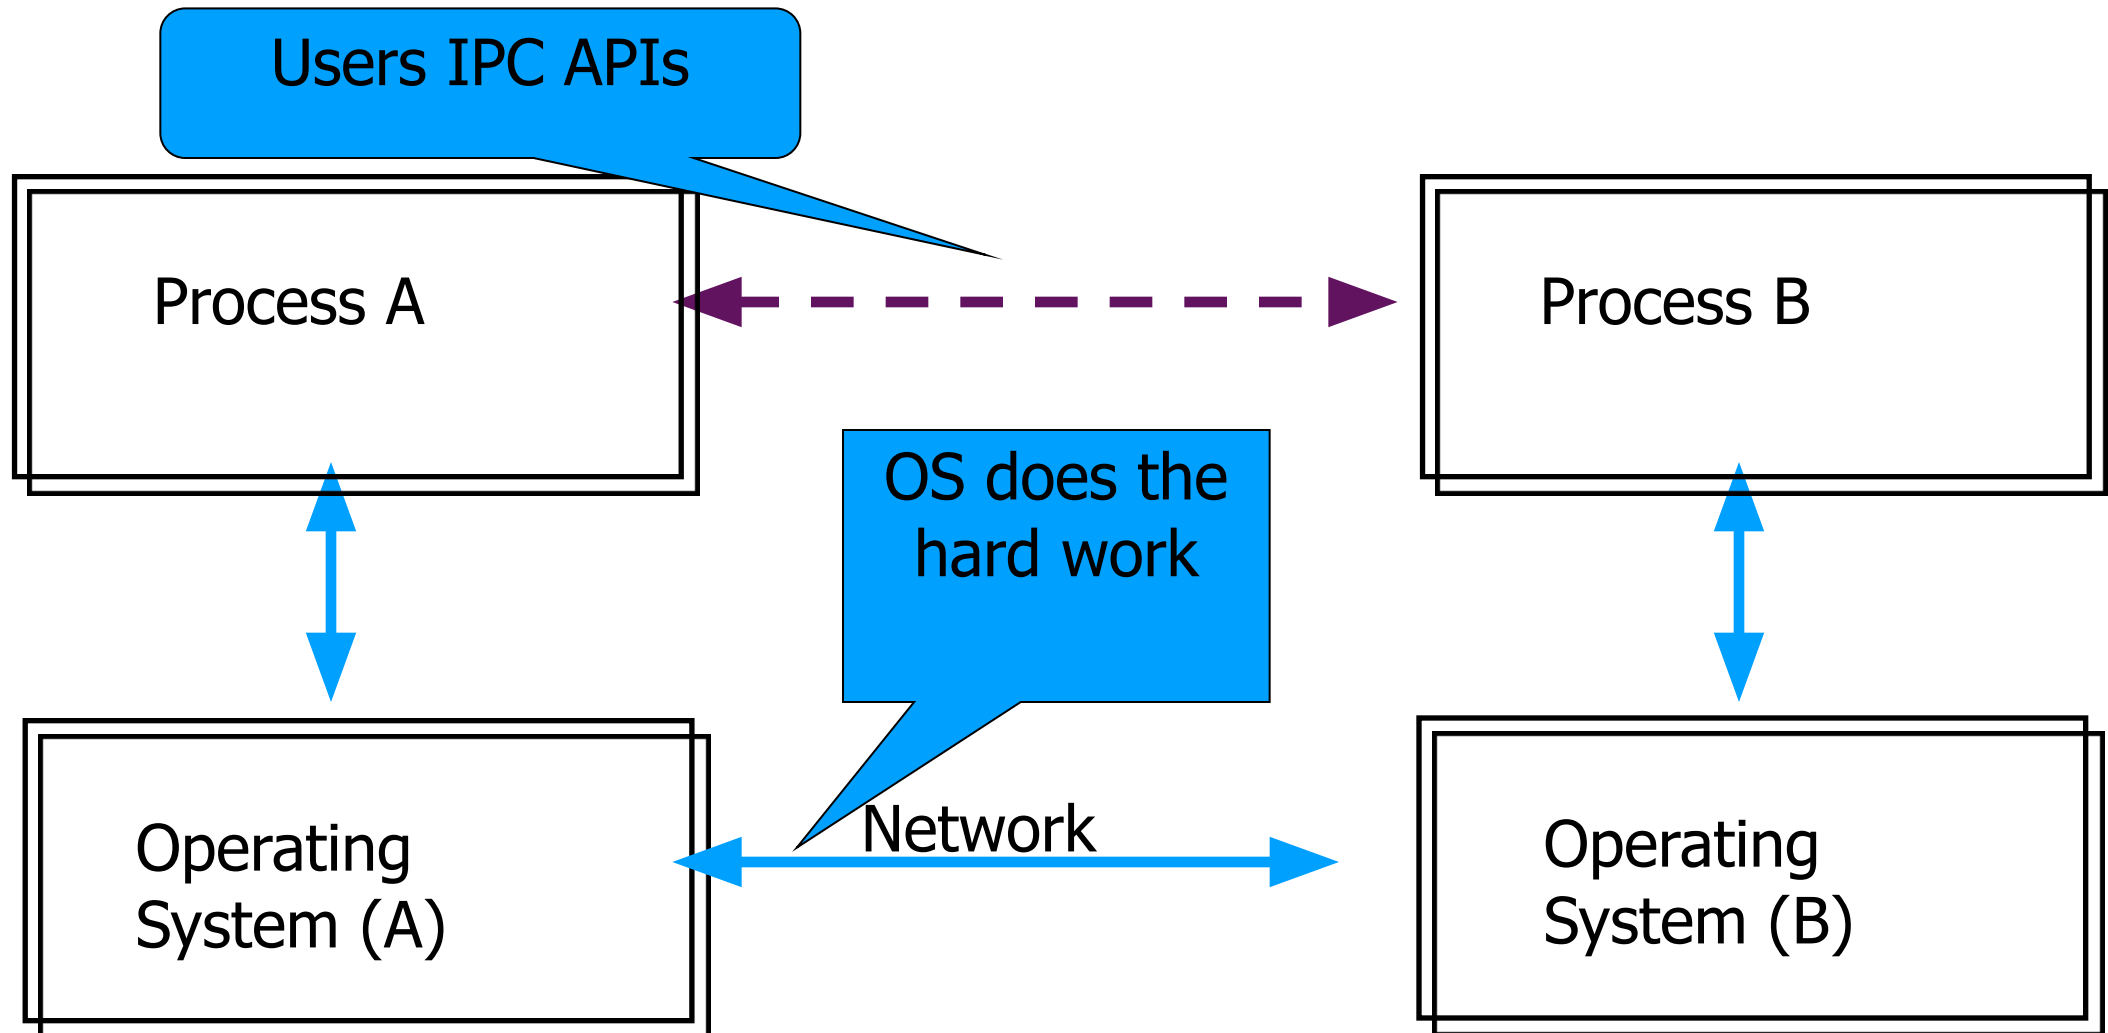
How Can Processes communicate?

- IPC mechanisms include:
 - Pipes and sockets (**stream** oriented).
 - Processes communicate in continuous *streams of bytes* sent over persistent connections.
 - Shared memory (**memory** oriented).
 - Processes communicate with each other through shared variables in *memory*.
 - Message passing (**message** oriented).
 - Processes communicate by sending each other discrete *chunks of data - messages*.
 - Remote procedure call (**procedure** oriented).
 - Processes communicate using *procedure call*

Single System IPC



IPC Across Different Systems



Towards Common APIs

- *Application Programming Interface* (API) is the generic name for an interface to some library of software functions. It is a connection between computers or between computer programs to allow communications.
- We first review some “traditional” APIs for IPC - specifically applicable to the “single-system” case.
- *But*, has become increasingly common to adopt the more general “different-systems” APIs, *even* when the communicating processes happen to run on the same system
 - In recent years much work has been put into developing these APIs for *distributed systems*.


SINGLE SYSTEM IPC

Traditional IPC: Pipes

- One of the simplest forms of single-system IPC
 - Still ubiquitous in UNIX-like systems (e.g. Linux).
- A UNIX *pipe* has an input and an output.
 - A stream of bytes is written to the output; the same stream of bytes is read from the input.
 - The inputting process will *block* if there is no data currently in the pipe.
- A pipe is created by some *parent process*, then used to communicate between (typically) two *child processes*.

Stream-based IPC Using Pipes


- **write** and **read** like a file, but more efficient.
 - One way byte stream.
 - Kernel buffers the data.
 - Operates as a *bounded buffer* with blocking (buffer is 64KB in Linux).



Scenario

- Most familiar use is in UNIX shell, e.g.:


```
$ ls | more
```



1. Shell creates pipe with two ends (**pipe** system call)
2. Shell creates two child processes (**fork** system calls), passing one end of pipe to each.
3. Child processes exec **ls** and **more** commands respectively.
4. They communicate through the inherited pipe.

Traditional IPC: Shared Files

- Require file locking or record locking to allow cooperating processes to share a resource safely.
 - File – locks the whole file;
 - Record – locks portion of a file.



Traditional IPC: System V

- *System V* was a dialect of UNIX developed in the 1980s.
- Many features adopted into “Portable Operating System Interface” (*POSIX*) standards, and still available today in Linux, etc.
- It incorporated *APIs* for single-system IPC supporting, for example:
 - Shared memory segments
 - Semaphores
 - Message queues

Shared Memory

- We have emphasized in earlier lectures that processes (in contrast to threads) have *private* address spaces, and in general *don't* share memory.
- This is the general rule, but specific *system calls* can be used to *create* memory areas that can be accessed by multiple processes.

System V Shared Memory API

- Creating shared memory segments is generally done in systems languages like C, e.g.:

```
key_t key;  
int shmid;  
char *data;  
  
key = ftok("/home/beej/somefile3", 'R');  
shmid = shmget(key, 1024, 0644 | IPC_CREAT);  
data = shmat(shmid, (void *)0, 0);
```

- Creates and attaches to a shared memory segment of size 1024 bytes, identified by **key**.

Memory Mapped Files

- On modern UNIX-like systems it is more common to implement shared memory segments by using the *virtual memory system* explicitly.
- A process can explicitly map a specified file into their *memory space* using the POSIX function **mmap**.
- If two or more processes map *the same file*, this effectively creates a shared memory region.

Shared Memory: Issues

- Once we have shared memory, the same issues of *data consistency* arise as with threads
 - How to prevent processes getting in each others way when performing critical updates?
 - How to make sure processes access shared data in some orderly patterns?
- The same kinds of solutions apply, and System V and POSIX provide *semaphores* that can be accessed from multiple processes.

IPC ACROSS COMPUTERS

IPC via Message Passing

- Processes interact by sending and receiving messages
 - Similar to pipes, but messages are isolated data chunks of specified size, rather than *unstructured streams* of bytes, as in pipes.
 - Sometimes say message passing is *connectionless*.
- As with pipes, problems of interference that arise when accessing shared data are avoided, *and the model works for communication between processes on different computers.*

Message Passing: Issues

- Problem: how many messages should be buffered temporarily during communication?
- Solutions:
 - **Zero-capacity queues:** 0 messages,
 - Sender *always waits for receiver* (synchronization is called *rendezvous*).
 - Sounds restrictive, but simplifies formal analysis of distributed programs (*CSP*, *CCS*, *Pi Calculus*, etc)
 - **Bounded capacity queues:** **finite** length of n messages,
 - Sender waits *if* link buffer is full (e.g. *MPI*).
 - **Unbounded capacity queues:** **infinite** queue length,
 - Sender *never waits*
 - Message queue in week 4 appendix, for example.

Message Passing: Implementations

- For general messaging there is the *Java Message Service* (JMS) API, implemented by various projects and vendors, e.g.:
 - Sun (Oracle) Java System Message Queue
 - BEA Weblogic
 - IBM WebSphere
- For *parallel computing* there is the *Message Passing Interface* (MPI), implemented by open source projects and hardware vendors.

Sockets

- *Sockets* provide a *programming model* with some features of message passing, though they are most commonly used for *stream-oriented* communication.
- In this respect they are similar to pipes. But unlike pipes, sockets can connect unrelated processes – including processes on *different* computers.

Socket API


- *Berkeley Sockets* API[†] implemented by *system calls* in Linux or Windows:

Primitive	Meaning
SOCKET	Create a new communication end point
BIND	Attach a local address to a socket
LISTEN	Announce willingness to accept connections; give queue size
ACCEPT	Block the caller until a connection attempt arrives
CONNECT	Actively attempt to establish a connection
SEND	Send some data over the connection
RECEIVE	Receive some data from the connection
CLOSE	Release the connection

Using Sockets

- After initializing a socket by `socket()`, `bind()`, `listen()`, “server” waits for connections on specified port by calling `accept()`.
- “Client” calls `connect()`, passing in a local socket and address of server socket, (IP address plus port number).
 - If connection succeeds, a new socket is returned by `accept()`.
 - Client and server exchange byte arrays of data over the socket pair using `send()` and `recv()` calls.

Scenario



1. Server creates socket (**socket** system call), binds it to an IP address and port (**bind**), and marks it as a server socket (**listen**)
2. Server does **accept**, to wait for connection from client.
3. Client creates socket (**socket**)...
4. ... and connects to IP address and port for server socket (**connect**).
5. A new socket is created on server connected to client socket; hosts communicate.
6. Server may do **accept** again, waiting for other connections...

Remote Procedure Call (RPC)

- Suggested by Birell and Nelson in 1984.
- Goals:
 - Access-transparent call semantics.
 - Remote calls *look like* local procedure calls.
 - Require to convert calls into network messages.
- Server:
 - Exports modules of procedures (somehow).
- Clients:
 - Call these procedures (somehow).
 - Look as close as possible to local-procedure call, *but* really network messages.

How RPC Works?

Client Node

```
void main()  
{  
  ...  
  i = y(a,b);  
  ...  
}
```

(a,b)

result

Server Node

```
int y(c,d) { ... }
```

(a,b)

result

System A.

System B

Why RPC ?

- Extends the conventional procedure call to the client/server model.
 - Remote procedures accept arguments and return results.
 - Makes it easy to design and understand programs.
 - Helps programmer to focus on the application instead of the communication protocol.
 - Allows a client to execute procedures on other computers.
 - Simplifies the task of writing client/server programs.
- ☐ RPC *forms the foundation for many distributed utilities* used today, like “Network File System” NFS and “Network Information Service” NIS in UNIX-based systems.

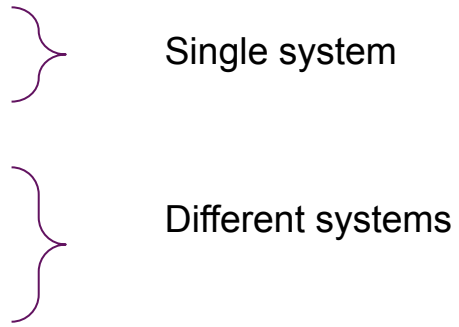
RPC: Issues

- Transparency:
 - Syntactic transparency: RPC should have *same syntax* as local procedure call.
 - Semantic transparency: RPC *semantics should be identical* to local procedure call.
- Standard representation: Need for external data representation (XDR) for all data types.
 - One machine may be *little-endian* and the other machine may be *big-endian*.
 - One machine may be using ASCII and the other UTF-16.
 - Representation of floating point numbers on the two machines may be different.

RPC: Implementations


- *Common Object Request Broker Architecture* (**CORBA**).
- Java *Remote Method Invocation* (**RMI**).
 - See next week's lab
- *SOAP* (formerly *Simple Object Access Protocol*) Web services.
- *Windows Communication Foundation* (**WCF**).

Summary

- IPC mechanisms:
 - Shared Memory (memory based).
 - Pipe and File (stream based).
 - Message Passing (message based).
 - Sockets (commonly stream based).
 - RPC (procedure based).
 - We've only scratched the surface of IPC.
 - *Next Lecture – File Systems*
- 
- The diagram consists of two purple curly braces on the right side of the list. The top brace groups 'Shared Memory (memory based)' and 'Pipe and File (stream based)' under the label 'Single system'. The bottom brace groups 'Message Passing (message based)', 'Sockets (commonly stream based)', and 'RPC (procedure based)' under the label 'Different systems'.

Further Reading

- Andrew S. Tanenbaum and David J. Wetherall, “*Computer Networks*”, 5th Edition, Pearson, 2013 (CN)
 - Sockets discussed section 6.1.3.



UNIVERSITY_{OF}
PORTSMOUTH

Questions?

